


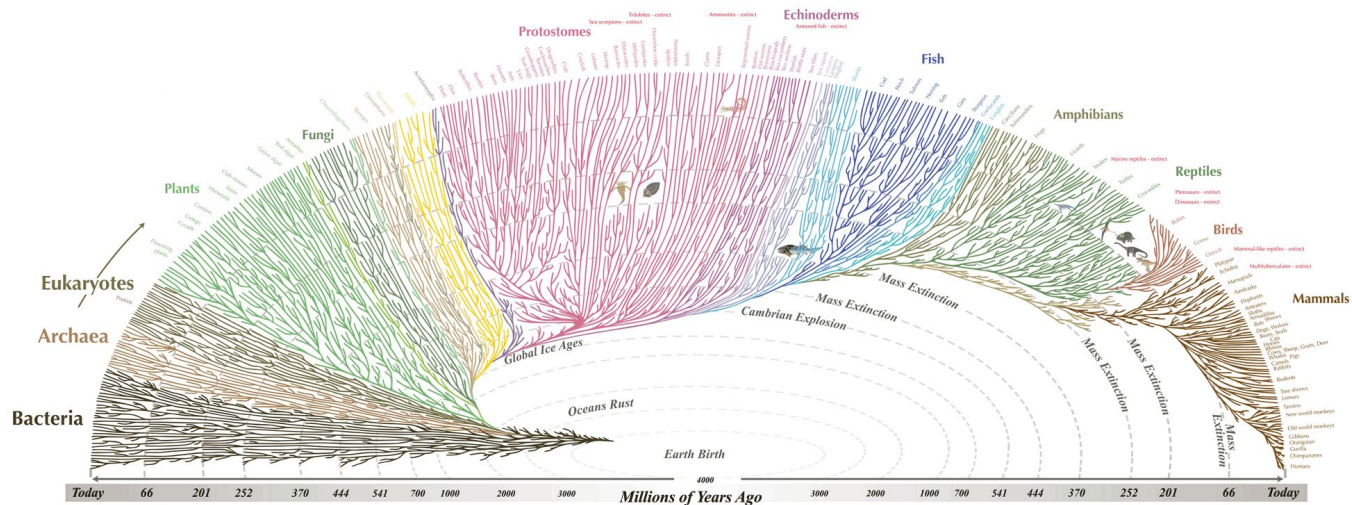
Week 2: Sequence alignment & search

- Sequence alignment problem
- Dynamic programming
- Global alignment
 - Needleman-Wunsch algorithm
- Local alignment
 - Smith-Waterman algorithm
- Substitution matrix
 - Construction & properties
- Fast sequence searches
 - BLAST; Statistics of similarity search

Sequence evolution



Then between A & B. various
 loss of relation. C & B. the
 first predation, B & D
 rather greater distance than
 then former would have
 formed. - heavy relation




All the major and many of the minor living branches of life are shown on this diagram, but only a few of those that have gone extinct are shown. Example: Dinosaurs - extinct

© 2008, 2017 Leonard Eisenberg. All rights reserved. evogram.com


deletion mutation insertion

ACATGGTCA → AC*TGGTCA → ACTGATCA → ACTGATICA

Evolutionary time



What is sequence alignment?



Sequences can be aligned by allowing for **gaps** and **mismatches**.

Alignment 1

ACTGATTCA

ACGCA-TCA

Alignment 2

ACTGATTCA

AC-GCATCA

Alignment 3

ACTG-ATTCA

AC-GCAT-CA

Which alignment is correct?

Alignment is gap placement.


Dynamic programming

Hemachandra/Fibonacci numbers: 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144,

$$\begin{aligned} F_0 &:= 0; F_1 := 1; \\ F_n &= F_{n-1} + F_{n-2}, \text{ for all } n \geq 2. \end{aligned}$$

A trivial algorithm for computing F_n :

```
naive_fib(n):  
    if n ≤ 1: return n  
    else: return naive_fib(n - 1) +  
           naive_fib(n - 2)
```



Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

Step 2

top
left
diagonal

	—	G	C	A	T
—					
G					
A					
T					

Substitution matrix to measure similarity in sequence alignments

Substitution matrix: A collection of scores for aligning nucleotides or amino acids with one another.

- Each score: the relative ease with which one nuc or AA may mutate into or substitute for another.
- Purely statistical, nothing directly to do with structure/biochemistry.

Ala	4																					
Arg	-1	5																				
Asn	-2	0	6																			
Asp	-2	-2	1	6																		
Cys	0	-3	-3	-3	9																	
Gln	-1	1	0	0	-3	5																
Glu	-1	0	0	2	-4	2	5															
Gly	0	-2	0	-1	-3	-2	-2	6														
His	-2	0	1	-1	-3	0	0	-2	8													
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4												
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4											
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5										
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5									
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6								
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7							
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4						
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5					
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11				
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7			
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4		
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val		

TITLE

CITED BY

YEAR


Basic local alignment search tool


SF Altschul, W Gish, W Miller, EW Myers, DJ Lipman

Journal of molecular biology 215 (3), 403-410

136003 *

1990

 U.S. National Library of Medicine

 National Center for Biotechnology Information

Sign in to NCBI

BLAST® Home Recent Results Saved Strategies Help


Basic Local Alignment Search Tool

BLAST finds regions of similarity between biological sequences. The program compares nucleotide or protein sequences to sequence databases and calculates the statistical significance.


[Learn more](#)

NEWS

IgBLAST 1.8.0 released
A new version of IgBLAST is now available.
Wed, 15 Nov 2017 16:00:00 EST

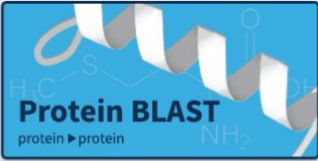
 More BLAST news...

Web BLAST

**Nucleotide BLAST**
nucleotide ► nucleotide

blastx
translated nucleotide ► protein

tblastn
protein ► translated nucleotide

**Protein BLAST**
protein ► protein

BLAST Genomes

Search

Human Mouse Rat Microbes

<https://www.ncbi.nlm.nih.gov/BLAST/>

What to brush-up on?

Biology

1. What is DNA? What does a DNA sequence look like? What do A, T, G, and C mean?
2. What is a protein? What does a protein sequence look like? What do the individual characters in the sequence mean?

Algorithms & coding

1. What is an algorithm?
2. What is a pseudocode of an algorithm?
3. What is recursion and what are loops (for, while)?
4. What is a conditional statement (if, else) and how is it used in coding?

What to brush-up on?

Analytical concepts & techniques

1. What is a matrix?
2. How do you write a mathematical expression to refer to a particular cell in the matrix based on its row and column?

Probability & statistics

1. What does probability mean?
2. How do you write a mathematical expression for the probability that: i) event **A** occurs, and ii) two events **A** and **B** occur together?
3. What is a probability distribution? What do the parameters in a probability distribution mean?
4. What is the difference between a discrete and a continuous probability distribution?
5. How do you write a mathematical expression for the probability that a particular variable **x** is less than or equal to a particular value **S**?
6. What is the binomial distribution? What kinds of processes does this distribution capture well?
7. What is the exponential distribution? What kinds of processes does this distribution capture well?

Week 2: Sequence alignment & search

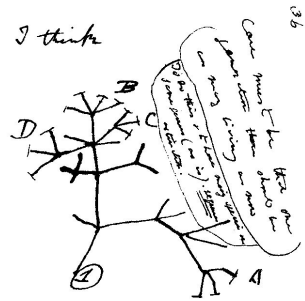
- Sequence alignment problem
- Dynamic programming
- Global alignment
 - Needleman-Wunsch algorithm
- Local alignment
 - Smith-Waterman algorithm
- Substitution matrix
 - Construction & properties
- Fast sequence searches
 - BLAST; Statistics of similarity search

Week 2: Sequence alignment & search

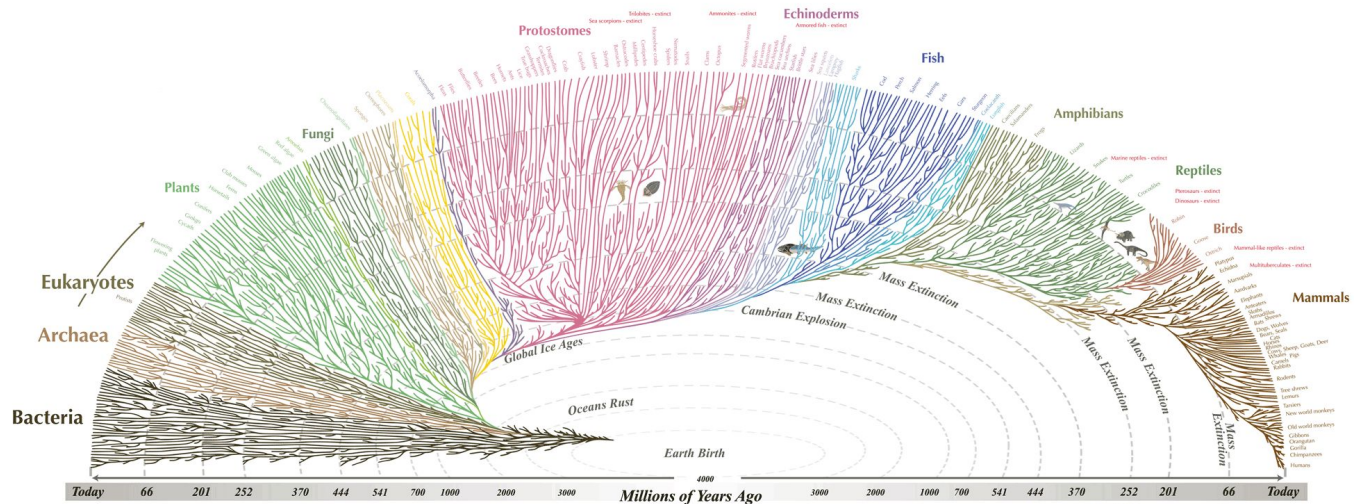
Alignment

- Sequence alignment problem
 - Dynamic programming
 - Global alignment
 - Needleman-Wunsch algorithm
 - Local alignment
 - Smith-Waterman algorithm
-
- Sign into Pear Deck using the link on Slack
 - Keep a paper and a pen(cil) ready

Sequence evolution



Then between A & B. various
 loss of relation. C & B. the
 first predation, B & D
 rather greater distance
 then former would have
 formed. - heavy relation




All the major and many of the minor living branches of life are shown on this diagram, but only a few of those that have gone extinct are shown. Example: Dinosaurs - extinct

© 2008, 2017 Leonard Eisenberg. All rights reserved.
 evoegen.com


deletion mutation insertion

ACATGGTCA → AC*TGGTCA → ACTGATCA → ACTGATICA

Evolutionary time



What is sequence alignment?



Sequences can be aligned by allowing for **gaps** and **mismatches**.

Alignment 1

ACTGATTCA
ACGCA-TCA

Alignment 2

ACTGATTCA
AC-GCATCA

Alignment 3

ACTG-ATTCA
AC-GCAT-CA

Which alignment is correct?

A scoring scheme:

- Match: **2**
- Mismatch: **-3**
- Gap: **-2**

We will come back to this!

$$2+2-3-3+2-2+2+2+2 \\ = 4$$

$$2+2-2+2-3-3+2+2+2 \\ = 4$$

$$2+2-2+2-2+2+2-2+2+2 \\ = 8$$

Alignment is gap placement.

How many possible alignments?

Dynamic programming

Solve a given complex problem by:

1. Breaking it into subproblems and
2. Storing the results of subproblems to avoid computing the same results again.

Two key properties of a problem that suggest that the given problem can be solved using DP.

1. Overlapping Subproblems
 - Given problem can be recursively broken down into subproblems that can be related to each other. That is, total no. of subproblems is polynomial.
2. Optimal Substructure
 - The optimal solution can be produced by combining optimal solutions of subproblems.



Richard Bellman

Optimal decision processes, involved time series & planning - thus 'dynamic' & 'programming'.

"It's impossible to use the word dynamic in a pejorative sense"; DP was "something not even a Congressman could object to."


Dynamic programming

Hemachandra/Fibonacci numbers: 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144,

$$\begin{aligned} F_0 &:= 0; F_1 := 1; \\ F_n &= F_{n-1} + F_{n-2}, \text{ for all } n \geq 2. \end{aligned}$$

A trivial algorithm for computing F_n :

```
naive_fib(n):  
    if n ≤ 1: return n  
    else: return naive_fib(n - 1) +  
           naive_fib(n - 2)
```



Dynamic programming


Hemachandra/Fibonacci numbers: $F_0 := 0$; $F_1 := 1$; $F_n = F_{n-1} + F_{n-2}$, for all $n \geq 2$.

Never recompute a subproblem $F(k)$, $k \leq n$, if it has been computed before.

Memoization: Remembering previously computed values.

Improved algorithm for computing F_n :

```
memo = { }  
  
fib(n):  
    if n in memo: return memo[n]  
    else if n = 0: return 0  
    else if n = 1: return 1  
    else: f = fib(n - 1) + fib(n - 2)  
    memo[n] = f  
    return f
```



These values are already computed and stored in memo when runtime processes these nodes of the recursion.

Dynamic programming

1. Overlapping subproblems
2. Optimal substructure

DP \approx recursion + memoization (reuse)

- Remember (memoize) previously solved “subproblems”; e.g., in Fibonacci, we memoized the solutions to the subproblems F_0, F_1, \dots, F_{n-1} , while unraveling the recursion.
- If we encounter a subproblem that has already been solved, reuse solution.
- Runtime \approx (no. of subproblems) * (time per subproblem)

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align **GCAT** with **GAT**

Step 1

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

	—	G	C	A	T
—					
G					
A					
T					

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align **G**CAT with **G**AT

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

Step 2

top
left
diagonal

	—	G	C	A	T
—					
G					
A					
T					

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align **G**CAT with **G**AT

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

Step 2

top
left
diagonal

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1				
A	-2				
T	-3				

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

Step 2

top
left
diagonal

	-	G	C	A	T
-	0	-1	-2	-3	-4
G	-1	?			
A	-2				
T	-3				

	-	G	C	A	T
-	0	-1	-2	-3	-4
G	-1	-2			
A	-2				
T	-3				

	-	G	C	A	T
-	0	-1	-2	-3	-4
G	-1	-2			
A	-2				
T	-3				

	-	G	C	A	T
-	0	-1	-2	-3	-4
G	-1	1			
A	-2				
T	-3				

	-	G	C	A	T
-	0	-1	-2	-3	-4
G	-1	1			
A	-2				
T	-3				

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align **GCAT** with **GAT**

Fill the remaining cells
in this matrix.

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

Step 2

top
left
diagonal

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1			
A	-2				
T	-3				

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align GCAT with GAT

p: gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

Step 2

top
left
diagonal

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1	0	-1	-2
A	-2	0	0	1	0
T	-3	-1	-2	0	2

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align GCAT with GAT

p: gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

top

left

diagonal

Step 3

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1	0	-1	-2
A	-2	0	0	1	0
T	-3	-1	-2	0	2

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align **G**CAT with **G**AT

What is the alignment?

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

top
left
diagonal

Step 3

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1	0	-1	-2
A	-2	0	0	1	0
T	-3	-1	-2	0	2

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align GCAT with GAT

GCAT
G-AT

p: gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

top
left
diagonal

Step 3

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1	0	-1	-2
A	-2	0	0	1	0
T	-3	-1	-2	0	2

Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

A scoring scheme:

- Match: **1**
- Mismatch: **-2**
- Gap: **-1**

Align **ATGCT** with **ATTACA**

p : gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = j * p; M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left(\begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array} \right)$$

top


left

diagonal


	-	A	T	T	A	C	A
-							
A							
T							
G							
C							
T							

Global & local alignment


A local alignment of strings s and t is an alignment of a substring of s with a substring of t .



Global Alignment



Local Alignment



Smith-Waterman algorithm

Similar to Needleman-Wunsch, with 3 changes:

- First row/column set to 0.
- No negative scores; set to 0. (Don't record direction.)
- Backtrack from cell with highest score & stop at 0.

p: gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = 0; M(i, 0) = 0$$

$$M(i, j) = \text{MAX}(\begin{array}{l} 0, \\ M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array})$$

top

left

diagonal

Align GCAT with GCT

What are the values in the first row and first column?

	-	G	C	A	T
-					
G					
C					
T					

Smith-Waterman algorithm

Similar to Needleman-Wunsch, with 3 changes:

- First row/column set to 0.
- No negative scores; set to 0. (Don't record direction.)
- Backtrack from cell with highest score & stop at 0.

p: gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = 0; M(i, 0) = 0$$

$$M(i, j) = \text{MAX}(\begin{array}{l} 0, \\ M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array})$$

top

left

diagonal

Align GCAT with GCT

Fill this matrix and
enter the highest value.

	-	G	C	A	T
-					
G					
C					
T					

Smith-Waterman algorithm

Similar to Needleman-Wunsch, with 3 changes:

- First row/column set to 0.
- No negative scores; set to 0. (Don't record direction.)
- Backtrack from cell with highest score & stop at 0.

p: gap penalty

$s(S1_i, S2_j)$: match/mismatch score

$$M(0, j) = 0; M(i, 0) = 0$$

$$M(i, j) = \text{MAX}(\begin{array}{l} 0, \\ M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + s(S1_i, S2_j) \end{array})$$

top

left

diagonal

Align GCAT with GCT

GC
GC

	-	G	C	A	T
-	0	0	0	0	0
G	0	1	0	0	0
C	0	0	2	1	0
T	0	0	1	1	2

Week 2: Sequence alignment & search

Substitution, BLAST

- Substitution matrix
 - Construction & properties
- Fast sequence searches
 - BLAST; Statistics of similarity search

Substitution matrix to measure similarity in sequence alignments



Dr. Margaret Dayhoff

Applying math & computational techniques to the sequencing of proteins and nucleic acids.

- 1965: First collection of protein seqs.
- Single-letter code for amino acids.
- 1966: 'Evolutionary trees'.
- **1978: First AA similarity-scoring matrix.**
- 1980: Launched the Protein Information Resource, the first online database system that could be accessed by telephone line.

Substitution matrix: A collection of scores for aligning nucleotides or amino acids with one another.

- The scores represent the relative ease with which one nucleotide or amino acid may mutate into or substitute for another.
- Purely statistical, nothing directly to do with structure/biochemistry.

[illegible]

Substitution matrix to measure similarity in sequence alignments

Substitution matrix: Each score is a log-odds score equal to the logarithm of the ratio of the likelihoods of two hypotheses: i) the residues can substitute for one another, or ii) not.

$$s(a,b) = \frac{1}{\lambda} \log \frac{p_{ab}}{f_a f_b}$$

- p_{ab} : likelihood of these two residues being correlated because they're homologous.
 - p_{ab} are the target frequencies: the probability that we expect to observe residues a and b aligned in homologous sequence alignments.
- $f_a f_b$: likelihood of these two residues being uncorrelated and unrelated, occurring independently.
 - f_a and f_b are background frequencies: the probabilities that we expect to observe amino acids a and b on average in any protein sequence.
- λ : a scaling factor, usually set to something that helps round off all the terms in the score matrix to sensible integers.

Substitution matrix to measure similarity in sequence alignments

Substitution matrix: Each score is a log-odds score equal to the logarithm of the ratio of the likelihoods of two hypotheses: i) the residues can substitute for one another, or ii) not.

$$s(a,b) = \frac{1}{\lambda} \log \frac{p_{ab}}{f_a f_b}$$

- p_{ab} : likelihood of these two residues being correlated because they're homologous.
- $f_a f_b$: likelihood of these two residues being uncorrelated and unrelated, occurring independently.
- λ : a scaling factor

Assuming that each aligned residue pair is statistically independent of the others (biologically dubious, but mathematically convenient):

- The score of an alignment ("**alignment score**") = sum of individual log-odds scores for each aligned residue pair.

Substitution matrix to measure similarity in sequence alignments

BLOSUM (BLOcks SUBstitution Matrix) for protein sequence alignment.

- Scan the BLOCKS database for very conserved regions of protein families (w/o gaps in the alignment).
- Count the relative frequencies of AA and their substitution probabilities.
- Calculate a log-odds score for each of the 210 possible substitution pairs of the 20 standard amino acids.

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

Substitution matrix to measure similarity in sequence alignments

BLOSUM (BLOcks SUBstitution Matrix) for protein sequence alignment.

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

- The rarer the amino acid is, the more surprising it would be to see two of them align together by chance.
- L/L more common than WW:
 - $p_{LL} = 0.0371$, $p_{WW} = 0.0065$
- W is a much rarer amino acid:
 - $f_L = 0.099$, $f_W = 0.013$.

Check with
 $\lambda = 0.347$.

Substitution matrix to measure similarity in sequence alignments

BLOSUM (BLOcks SUBstitution Matrix) for protein sequence alignment.

- A/L pairs are slightly more frequent in homologous alignments than K/E pairs:
 - $p_{AL} = 0.0044$, $p_{KE} = 0.0041$.
- But, A and L are more common amino acids:
 - $f_A = 0.074$, $f_L = 0.099$, $f_K = 0.058$, $f_E = 0.054$.

Check with
 $\lambda = 0.347$.

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

Substitution matrix to measure similarity in sequence alignments

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

BLOSUM_r: the matrix built from blocks with less than **r**% of similarity.

- E.g., BLOSUM62: built using sequences with less than 62% similarity (sequences with $\geq 62\%$ identity were clustered)
 - BLOSUM62 among the best for detecting most weak protein similarities.
 - Default matrix for protein BLAST.

Substitution matrix to measure similarity in sequence alignments

Substitution
matrix for DNA

A				
T				
G				
C				
	A	T	G	C

Making-up an arbitrary matrix by fixing the p_{ab} values \rightarrow directly describes what homologous alignments are expected to look like.

- The resulting score matrix is optimal for detecting alignments that match these target frequencies.

Say, the matrix should be optimized for finding 88% identity alignments.

- Assume that all mismatches are equiprobable, and composition of both alignments and background sequences is uniform at 25% for each nucleotide ($f_a, f_b = 0.25$ for all a, b). Then,
 - Four identities: $p_{aa} = 0.22$
 - 12 types of mismatch: $p_{ab} = 0.01$.
- If we set $\lambda = 1$, this gives +1.26 for a match and -1.83 for a mismatch.
- Setting $\lambda = 0.25$ and round off: we have a new scoring system of +4/-7.

Substitution matrix to measure similarity in sequence alignments

Substitution
matrix for DNA

A				
T				
G				
C				
	A	T	G	C

Given a scoring matrix, we can back calculate target frequencies if two conditions are met:

$$s(a,b) = \frac{1}{\lambda} \log \frac{p_{ab}}{f_a f_b}$$

1. It must have at least one positive score, and
2. The expected score for random sequence alignments must be negative.

True for most score matrices:

- These properties are necessary to make local sequence alignment algorithms like BLAST and Smith-Waterman work.
- Both conditions are met by definition for matrices derived as log-odds scores, except for the useless case of $p_{ab} = f_a f_b$ for all a,b.

Examples:

- FASTA & WU-BLASTN: arbitrary +5/−4 scoring system; Optimal for detecting alignments that are 65% identical.
- NCBI BLASTN: +1/−2 scoring system; Optimal for detecting alignments that are 95% identical.

How do we scale this up to search an entire sequence database?

Given a query sequence, and a large set of target sequences (millions), which target sequences (if any) are related to the query?

- Individual alignments need not be perfect: Once initial matches are found, they can fine-tune them later.
- Must be very fast.

Exploit the nature of the problem (most sequences will be unrelated to the query):

- If any match with % identity ≤ 90 is going to be rejected, can ignore sequences which don't have a stretch of 10 nucleotides in a row.
- Pre-screen sequences for common long stretches.
- Pre-process the database offline and index k-mers.

TITLE

CITED BY

YEAR


Basic local alignment search tool

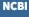
SF Altschul, W Gish, W Miller, EW Myers, DJ Lipman

Journal of molecular biology 215 (3), 403-410

136003 *

1990

 U.S. National Library of Medicine

 National Center for Biotechnology Information

Sign in to NCBI

BLAST® Home Recent Results Saved Strategies Help


Basic Local Alignment Search Tool

BLAST finds regions of similarity between biological sequences. The program compares nucleotide or protein sequences to sequence databases and calculates the statistical significance.


[Learn more](#)

NEWS

IgBLAST 1.8.0 released
A new version of IgBLAST is now available.
Wed, 15 Nov 2017 16:00:00 EST


 More BLAST news...

Web BLAST

**Nucleotide BLAST**
nucleotide ► nucleotide

blastx
translated nucleotide ► protein

tblastn
protein ► translated nucleotide

**Protein BLAST**
protein ► protein

BLAST Genomes

Search

Human Mouse Rat Microbes

<https://www.ncbi.nlm.nih.gov/BLAST/>

BLAST

Query sequence: PQGEFG

Word 1: PQG

Word 2: QGE

Word 3: GEF

Word 4: EFG

query word ($W = 3$)

Query: GSVEDTTGSQSLAALLNKCKT**PQG**QRLVNQWIKQPLMDKNRIEERLNLVEAFVEDAELRQTLQEDL

neighborhood
words

PQG	18
PEG	15
PRG	14
PKG	14
PNG	13
PDG	13
PHG	13
PMG	13
PSG	13
PQA	12
PQN	12
etc...	

neighborhood
score threshold
($T = 13$)


Query: 325 SLAALLNKCKT**PQG**QRLVNQWIKQPLMDKNRIEERLNLVEA 365
+LA++L+ TP G R++ +U+ P+ D + ER + A
Sbjct: 290 TLASVLDCTVT**PMG**SRMLKRWLHMPVRDTRVLLERQQTIGA 330

High-scoring Segment Pair (HSP)

Some uses of BLAST

- Finding the right/relevant species:
 - If you have a DNA sequence from unknown species, BLAST can help identify the correct/related species.
- Finding protein domains:
 - If you a protein sequence (or a translated nucleotide sequence), BLAST can be used to look for known protein domains in the query sequence.
- Mapping the phylogeny of a gene/protein:
 - BLAST can be used to find potential homologs of your gene/protein of interest across many species, which you can then use to generate a phylogenetic tree.
- Mapping DNA to a known chromosome:
 - If you are sequencing a gene from a known species but have no idea of the chromosome location, BLAST can help you. BLAST will show you the position of the query sequence in relation to the hit sequences.
- Annotations:
 - BLAST can also be used to map gene/protein annotations from one organism to another.

Statistics of similarity search



Distribution of real (squares) & expected similarity scores (Gumbel extreme value distribution).

P-value:

- The probability of observing a score equal to or greater than the observed score S .

E-value:

- The expected number of HSPs with score at least S .
- $E = Kmne^{-\lambda S}$

Database E-value:

- E-value after thousands/millions of searches $\approx E * D$.

Bit score:

- Normalized raw score.